



UNIVERSITY OF CHITTAGONG

Department of Computer Science and Engineering

Session: 2021–2022

7th Semester

Final Project : Design and Implementation of a
MiniLang Compiler
Course Title : Compiler Lab
Course Code No. : CSE 712
Program : B.Sc Engineering

Submitted To:

Samia Muntaha

Lecturer

Department of Computer Science and Engineering
University of Chittagong

Submitted By:

Ahsanul Hoque

ID: 22701048

Yousha Shamim Khan

ID: 22701007

Department of Computer Science and Engineering

Date of submission: July 20, 2026

Contents

1. Language Design Decisions	2
1.1 Overview of MiniLang Features Implemented	2
1.2 Rationale for Simplifying Assumptions	2
2. Lexical Analysis	3
2.1 Regular Expressions and Token Definitions	3
2.2 Handling of Whitespace, Comments, and Errors	3
2.3 Complexity Analysis: Time and Space Complexity of Scanning	3
3. Syntax Analysis	4
3.1 Grammar Design and Precedence Rules	4
3.2 Conflict Resolution Strategy	4
3.3 Complexity Analysis: Cost of Parsing	5
4. Abstract Syntax Tree (AST)	5
4.1 AST Structure and Node Types	5
4.2 Why an AST Instead of a Parse Tree	5
4.3 Complexity Analysis: Memory and Construction Cost	6
5. Symbol Table and Semantic Analysis	6
5.1 Scope Management Approach	6
5.2 Type-Checking Implementation	6
5.3 Error Detection Strategy	6
5.4 Complexity Analysis: Lookup, Insertion, and Traversal Cost	7
5.5 Robustness: Bounded Expression-Nesting Depth	7
6. Intermediate Code Generation	8
6.1 TAC Design and Instruction Format	8
6.2 Translation of Conditionals and Loops	8
6.3 Scope-Correct Naming (Shadowing)	8
6.4 TAC-Level Optimization	9
6.5 Complexity Analysis: Cost of AST Traversal and Code Emission	9
6.6 Empirical Validation	10
7. Overall Compiler Complexity	11
7.1 Combined Time Complexity of All Phases	11
7.2 Space Usage of AST, Symbol Table, and TAC	12
7.3 Discussion: Does the Compiler Operate in Linear Time?	12
8. Sample Compilation Results	12
8.1 Example MiniLang Program (<code>testcases/valid1_features.ml</code>)	12
8.2 Generated Three-Address Code (unoptimized)	13
8.3 Optimized Three-Address Code (written to <code>output.tac</code>)	14
8.4 Generated Target (Stack-Machine) Code (written to <code>output.asm</code>)	14
8.5 Verified Error Handling	14
8.6 Verified Scope Shadowing (<code>testcases/valid2_scoping.ml</code>)	15

A Flex/Bison compiler for MiniLang that performs lexical, syntactic, and semantic analysis, builds

an AST, generates Three-Address Code (TAC), applies TAC-level optimizations, and emits simplified stack-machine target code.

1. Language Design Decisions

1.1 Overview of MiniLang Features Implemented

MiniLang, as implemented, supports exactly the constructs required by the assignment specification:

- Variable declarations for two scalar types, `int` and `bool`, with an optional initializer: `int x;` or `int x = 5;`
- Assignment statements: `x = expr;`
- Arithmetic expressions with `+` `-` `*` `/`, standard precedence, and unary minus
- Relational expressions with `<` `>` `==` `!=`, producing a `bool` result
- `if/else` conditionals, including dangling-else resolution
- `while` loops
- `print(expr)` statements
- Block scoping with `{ }`, including shadowing of an outer variable by an inner declaration of the same name

1.2 Rationale for Simplifying Assumptions

- **Boolean literals.** The spec lists keywords `int`, `bool`, `if`, `else`, `while`, `print` but a usable `bool` type needs literal values, so `true` and `false` were added as keywords producing `BOOL_CONST` tokens. Without them, no MiniLang program could ever construct a boolean value directly.
- **No implicit type conversion.** Assignment and initialization require the expression type to exactly match the declared type (no `int/bool` coercion). This keeps the type system simple and makes the required “type mismatch” semantic check meaningful and easy to trigger deterministically in test cases.
- **Approximate (not per-token) source locations.** Errors are reported using the lexer’s current `yylineno` at the point a grammar rule reduces, rather than full `%locations` tracking with begin/end positions for every token. For a language with no multi-line expressions of real complexity, this gives accurate-enough line numbers while keeping the grammar actions simple.
- **Deterministic Makefile instead of the literal `*.c glob`.** The assignment’s example build command, `gcc lex.yy.c parser.tab.c *.c -o minicompiler`, silently double-compiles `lex.yy.c` and `parser.tab.c` on a normal shell, because the `*.c` glob matches those two generated files a second time after they are already named explicitly on the command line — this produces duplicate-symbol link errors (duplicate `main`, `yylex`, etc.) on every modern GCC. The provided Makefile therefore lists every source file explicitly while still using exactly the three required tools (`flex`, `bison -d`, `gcc`) in the same roles.
- **Reserved `$`-prefixed internal names.** Compiler-generated temporaries (`$t1`, `$t2`, ...) and labels (`$L1`, `$L2`, ...) use a leading `$`, a character the lexer’s identifier rule never accepts. This guarantees by construction that generated TAC names can never collide with a user-declared variable, without needing a runtime collision check.

2. Lexical Analysis

Implemented in `lexer.1` using Flex, with `%option yylineno` for automatic line tracking.

2.1 Regular Expressions and Token Definitions

```
DIGIT  [0-9]
LETTER [a-zA-Z_]
```

```
"int"          { return INT_KW; }
"bool"         { return BOOL_KW; }
"if" | "else" | "while" | "print"  { return <the matching keyword token>; }
"true" | "false"      { yylval.ival = 1 or 0; return BOOL_CONST; }
{LETTER}{(LETTER){DIGIT})* { yylval.sval = strdup(yytext); return IDENT; }
{DIGIT}+            { yylval.ival = atoi(yytext); return INT_CONST; }
"==" | "!=" | "<" | ">" | "="  -> EQ, NE, LT, GT, ASSIGN
"+" | "-" | "*" | "/"      -> PLUS, MINUS, STAR, SLASH
"(" | ")" | "{" | "}" | ";" | "," -> LPAREN, RPAREN, LBRACE, RBRACE, SEMI, COMMA
```

Keyword rules are listed before the general identifier rule. Flex resolves ties in matched length by rule order, so `"int"` and the identifier pattern both match the text `int` with equal length, but the keyword rule wins because it appears first — this is what lets keywords be recognized without a separate keyword-lookup table.

2.2 Handling of Whitespace, Comments, and Errors

- Whitespace (`[ \t\r]`) is discarded; newlines are matched separately so `yylineno` advances correctly.
- Line comments `// ...` are matched and discarded in one rule.
- Block comments `/* ... */` use a Flex start condition (`%x COMMENT`) so they can safely span multiple lines; an `<<EOF>>` rule reports an “unterminated comment” lexical error if the file ends before the comment is closed.
- Any other unrecognized character is caught by a catch-all `.` rule, which prints “Lexical Error at line N: unexpected character ‘...’” and increments a global `lexErrorCount`, then keeps scanning so multiple lexical errors in one file are all reported in a single pass instead of stopping at the first one.

2.3 Complexity Analysis: Time and Space Complexity of Scanning

Flex compiles all rules into a single deterministic finite automaton (DFA), so each input character is examined by a constant number of state transitions regardless of how many token rules exist. For an input of length n characters, scanning the entire source runs in $O(n)$ time. Because MiniLang has no unbounded-lookahead constructs (no arbitrarily long lexemes requiring backtracking across the whole file), the DFA never needs to rescan already-consumed input. Space usage is $O(1)$ with respect to input size for the automaton itself (its state table is fixed at generation time); the only per-token space is the buffer for the currently matched lexeme, plus one `strdup`’d string per identifier token, which is $O(1)$ additional space per token and is freed by the parser action that consumes it.

3. Syntax Analysis

Implemented in `parser.y` using Bison in LALR(1) mode, with `%define parse.error verbose` for descriptive syntax error messages.

3.1 Grammar Design and Precedence Rules

```
%left EQ NE
%left LT GT
%left PLUS MINUS
%left STAR SLASH
%right UMINUS
%nonassoc IFX
%nonassoc ELSE

program      : stmt_list
stmt_list    : /* empty */ | stmt stmt_list
stmt         : declaration | assignment | if_stmt | while_stmt | print_stmt | block
declaration  : type IDENT SEMI | type IDENT ASSIGN expr SEMI
if_stmt      : IF LPAREN expr RPAREN stmt %prec IFX
              | IF LPAREN expr RPAREN stmt ELSE stmt
while_stmt   : WHILE LPAREN expr RPAREN stmt
block        : LBRACE stmt_list RBRACE
expr         : expr (PLUS|MINUS|STAR|SLASH|LT|GT|EQ|NE) expr
              | MINUS expr %prec UMINUS
              | LPAREN expr RPAREN | IDENT | INT_CONST | BOOL_CONST
```

Operator precedence and associativity are resolved declaratively with Bison's precedence table rather than by restructuring the grammar into a precedence-climbing cascade of nonterminals (e.g. `expr -> term -> factor`). This keeps the grammar itself small (a single `expr` nonterminal) while still deterministically resolving every arithmetic/relational ambiguity: multiplication/division bind tighter than addition/subtraction, which bind tighter than relational operators, which bind tighter than equality operators, and unary minus binds tightest of all.

3.2 Conflict Resolution Strategy

Two categories of shift/reduce ambiguity are resolved explicitly:

- **Operator ambiguity** (e.g. is `a - b - c` `(a-b)-c` or `a-(b-c)`?) is resolved by the `%left/%right` precedence declarations above.
- **The dangling-else problem** (does `else` in `if (a) if (b) s1 else s2` bind to the inner or outer `if`?) is resolved with the classic Bison idiom: a pseudo-token `IFX` is given lower precedence than `ELSE`, and the `if-without-else` rule is tagged `%prec IFX`. Because `ELSE` has the higher precedence, Bison prefers to shift the `else` token rather than reduce the shorter rule, which binds `else` to the nearest unmatched `if` — the conventional and expected behavior. Building with `bison -d parser.y` reports **zero unresolved conflicts**.

3.3 Complexity Analysis: Cost of Parsing

Bison generates a table-driven LALR(1) shift-reduce parser. Each input token triggers exactly one shift and, amortized, at most a small constant number of reduce actions (bounded by the grammar’s maximum rule length, which is fixed at compile-generation time and independent of input size). Consequently, parsing runs in $O(n)$ time for a token stream of length n . The parser stack grows with nesting depth (e.g. nested blocks, nested parenthesized expressions), so space usage is $O(d)$ where d is the maximum nesting depth of the program — in the worst case $O(n)$ for a pathologically deep single expression, but $O(1)$ additional space per token in the common case of shallow nesting. Bison’s default stack-growth ceiling (`YYMAXDEPTH`, ~10,000 by default) additionally guarantees that a pathologically deeply *parenthesized* input fails gracefully with a “memory exhausted” parse error rather than exhausting process memory — verified empirically in §6.6.

A subtlety that matters for this bound to actually hold: `stmt_list` is deliberately written as a *left*-recursive rule (`stmt_list: /* empty */ | stmt_list stmt`), not the more naturally-written right-recursive `stmt_list: /* empty */ | stmt stmt_list`. A right-recursive list forces an LALR parser to shift and hold every element on the parser stack until the final empty-production reduce, then cascade-reduce from the bottom up — meaning a program with n purely *sequential*, non-nested statements would need $O(n)$ parser-stack depth, not $O(d)$. That would make d in the paragraph above silently mean “statement count” for a wide program, not “nesting depth” as claimed, and a plain, non-adversarial MiniLang program of a few thousand statements would fail to parse for a reason with nothing to do with nesting. Left recursion lets Bison reduce incrementally ($O(1)$ stack per statement) at the cost of building each statement list back-to-front; `reverseStmtList()` (`ast.c`) restores source order in one $O(n)$ pass wherever a `stmt_list` is finalized (`program` and `block`). This keeps d accurately meaning nesting depth, confirmed empirically in §6.6 (Experiment B).

4. Abstract Syntax Tree (AST)

4.1 AST Structure and Node Types

`ast.h` defines a single tagged-union node type (`ASTNode`) discriminated by a `NodeType` enum, with a C union holding only the fields relevant to that node kind:

- **Leaves:** `NODE_INT_LIT`, `NODE_BOOL_LIT`, `NODE_IDENT`
- **Expressions:** `NODE_BINOP` (op, left, right), `NODE_UNARY_MINUS` (expr)
- **Statements:** `NODE_VAR_DECL` (type, name, optional init), `NODE_ASSIGN` (name, expr), `NODE_IF` (cond, then, optional else), `NODE_WHILE` (cond, body), `NODE_PRINT` (expr), `NODE_BLOCK` (stmt list)
- **Sequencing:** `NODE_STMT_LIST`, a singly-linked list node (stmt, next) used to chain statements within a block or at the top level

Every node also carries a source line number, propagated from `yylineno` at the point the corresponding grammar rule reduced, so that later phases (semantic analysis) can report errors against the original source line.

4.2 Why an AST Instead of a Parse Tree

A full concrete parse tree would retain a node for every grammar symbol matched, including punctuation tokens (`( ) { } ;`) and the intermediate nonterminals introduced purely to encode

precedence. None of that is needed by semantic analysis or code generation — a `BINOP` node only needs its operator and two operand subtrees, not a matched pair of parenthesis tokens. Building an AST directly in the parser’s semantic actions (each grammar rule constructs exactly one AST node from its already-reduced children) discards this syntactic noise at the point of creation, keeping the tree proportional to the program’s actual structure rather than to its concrete grammar derivation.

4.3 Complexity Analysis: Memory and Construction Cost

Exactly one `ASTNode` is allocated per reduced grammar rule that produces a node (one per literal, identifier use, operator application, and statement). For a program whose token stream has length n , the number of AST nodes is $O(n)$, and since each node has a fixed number of fields (bounded by the largest union member), total AST memory is $O(n)$. Construction cost is $O(1)$ per node (a single `malloc` and field assignment), so building the whole tree is $O(n)$ overall, done incrementally as the parser reduces — there is no separate tree-construction pass.

5. Symbol Table and Semantic Analysis

5.1 Scope Management Approach

`symbol_table.c` implements scopes as a linked stack of hash tables. Each `Scope` owns a hash table (separate chaining, djb2 string hash), starting at 211 buckets, and a parent pointer to the enclosing scope. `enterScope()/exitScope()` push and pop a scope; `declareSymbol()` inserts into (and rejects a duplicate within) only the current scope’s table; `lookupSymbol()` walks from the current scope outward to the global scope, returning the innermost match — which is exactly what makes an inner declaration correctly shadow an outer one of the same name. A scope’s bucket count is not fixed for the scope’s lifetime: `declareSymbol` doubles it and rehashes every existing entry whenever the load factor (entries $\div$ buckets) would otherwise exceed 2 (§5.4 explains why this matters).

5.2 Type-Checking Implementation

`semantic.c` performs a single recursive-descent walk over the AST (`checkStmt / checkExpr`), threading a `ValueType` (`TYPE_INT`, `TYPE_BOOL`, or `TYPE_UNKNOWN` for an already-erroneous subexpression, which suppresses cascading errors) bottom-up through every expression:

- Arithmetic operators (+ - * /) require both operands `int` and produce `int`.
- Order operators (< >) require both operands `int` and produce `bool`.
- Equality operators (== !=) require both operands to have the same type (either both `int` or both `bool`) and produce `bool`.
- Unary minus requires an `int` operand.

5.3 Error Detection Strategy

All four required semantic checks are implemented directly against the scoped symbol table:

- **Undeclared variables** — any identifier use where `lookupSymbol` returns `NULL`.
- **Multiple declarations in the same scope** — `declareSymbol` returns 0 when the name already exists in the current scope’s bucket chain (an outer-scope match is a legal shadow, not an error).
- **Type mismatch in expressions** — checked at every binary operator, assignment, and initializer, as described above.

- **Invalid conditional expressions** — the condition of both `if` and `while` must type-check to `bool`.

Errors do not abort the walk: each is printed immediately with its source line and the walk continues, so a single compilation reports every semantic error found in one pass rather than only the first. Code generation only proceeds once the error count is zero.

5.4 Complexity Analysis: Lookup, Insertion, and Traversal Cost

With the djb2 hash distributing reasonably uniformly, `declareSymbol` (insertion) is $O(1)$ amortized within one scope, and `lookupSymbol` is $O(1)$ amortized per scope examined. This specifically requires the load-factor-triggered doubling described in §5.1: a scope’s bucket count *must* grow with its entry count, not stay fixed, or the average chain length — and therefore the true per-operation cost — grows linearly with however many names end up declared in that one scope, degrading toward $O(v)$ per operation for a scope with v entries instead of the $O(1)$ being claimed. (A fixed-size table is a reasonable choice only if v is known in advance to stay small; MiniLang has no such guarantee — a single, unnested sequence of declarations can be arbitrarily long — so a fixed bound understates the actual worst case rather than merely simplifying it. §6.6, Experiment B, is precisely this scenario and is why the growable version is not just a hygiene improvement.) With growth accounted for, `lookupSymbol` in the worst case examines every enclosing scope, giving $O(d)$ where d is the current scope-nesting depth — in practice a small constant for MiniLang programs, since d is bounded by how deeply `{ }` blocks are nested, independently of how many names live in any one of those scopes. The semantic walk itself visits every AST node exactly once, performing $O(1)$ amortized symbol-table work per node, so full semantic analysis of a program with n AST nodes runs in $O(n \cdot d)$ time, which is effectively $O(n)$ for realistic (shallow) nesting depths. Space usage is $O(v)$ where v is the number of live variable declarations across all currently-open scopes (each popped scope’s symbols, and its bucket array, are freed immediately in `exitScope`).

5.5 Robustness: Bounded Expression-Nesting Depth

Bison’s own parser-stack ceiling (§3.3) protects the compiler against pathologically *nested* input such as thousands of unmatched parentheses — parsing such a program fails cleanly with a syntax error before any AST is even built. It does **not**, however, protect against a pathologically *long* input of a different shape: a flat, left-associative chain of binary operators (e.g. `x = 1+1+1+...+1` with tens of thousands of terms). Because `+-/*//` are left-associative, Bison reduces each operator as soon as the next one of equal precedence is seen, so a long chain never grows the *parser’s* stack — but it still builds a left-deep AST of depth proportional to the chain length, and `checkExpr` (§5.2) walks that tree with one C function-call frame per nesting level. An unbounded chain would therefore exhaust the C call stack and crash the compiler with a segmentation fault rather than reporting a compile error.

`checkExpr` guards against this with an explicit depth counter, `MAX_EXPR_DEPTH = 3000` — chosen with comfortable headroom below both Bison’s own $\sim 10,000$ -level ceiling and the actual C-stack crash threshold observed empirically to be well above 32,000 levels (§6.6). Exceeding the limit reports a single, clean **Semantic Error: expression nested too deeply (limit 3000 levels)** and unwinds without recursing further, rather than allowing the recursion to continue toward a crash. A one-shot flag prevents the message from being reported more than once per compilation, since the global depth counter can cross the threshold multiple times while unwinding a deep left spine and then visiting shallow sibling subexpressions.

`genExpr` (§6.1) recurses over the same left-deep expression shape while emitting TAC, so it carries an independent, identically-configured `MAX_GEN_EXPR_DEPTH` guard as defense in depth. In the normal pipeline this second guard is unreachable — `main.c` runs semantic analysis first and halts before code generation if `semErrorCount > 0`, so a program that trips `checkExpr`’s limit never reaches `genExpr` at all — but it means `genExpr` does not silently assume a well-formed input and stays safe to call on its own.

6. Intermediate Code Generation

6.1 TAC Design and Instruction Format

`codegen.c` represents TAC as a dynamic array of a single `TACInstr` struct with a kind tag and `op / arg1 / arg2 / result` string fields, covering seven instruction kinds:

```
result = arg1           (TAC_ASSIGN, e.g. x = t1)
result = arg1 op arg2   (TAC_BINOP,  e.g. t1 = a + b)
result = op arg1        (TAC_UNOP,   e.g. t2 = -t1)
result:                 (TAC_LABEL,  e.g. $L1:)
goto result             (TAC_GOTO)
ifz arg1 goto result    (TAC_IFZ)
print arg1              (TAC_PRINT)
```

Fresh temporaries (`$t1`, `$t2`, ...) and labels (`$L1`, `$L2`, ...) are minted by simple monotonically-increasing counters. Every temporary produced while translating a subexpression is, by construction, consumed exactly once by its immediate parent — a property the optimizer’s dead-code pass relies on.

6.2 Translation of Conditionals and Loops

```
if (cond) S1 else S2:           while (cond) S:
  <code to eval cond -> c>      L_start:
  ifz c goto L_false           <code to eval cond -> c>
  <code for S1>                 ifz c goto L_end
  goto L_end                   <code for S>
  L_false:                     goto L_start
  <code for S2>                 L_end:
  L_end:
```

Both forms use exactly the textbook back-patched-label scheme, generalized so the if-without-else case simply omits the `goto L_end` / second block and reuses `L_false` as the join point.

6.3 Scope-Correct Naming (Shadowing)

Block scoping means two distinct declarations can share the same source name (an inner `int x` shadowing an outer one). TAC itself has no concept of scope, so `codegen.c` mirrors `semantic.c`’s scope stack with its own lightweight declare/resolve mechanism, `cgDeclare/cgResolve`: the first declaration of a name keeps its original name as its TAC storage location; a shadowing declaration is assigned a disambiguated name (`x$2`, `x$3`, ...). Every identifier use resolves through this table to the innermost visible declaration’s TAC name, exactly matching `lookupSymbol`’s resolution order, so the generated code observes the same shadowing behavior the semantic analyzer already

validated. A shadowed initializer is evaluated *before* the new binding is registered (mirroring §5.2’s own evaluation order), so a self-shadowing initializer such as `int x; { int x = x + 1; }` correctly reads the outer `x` on the right-hand side rather than the not-yet-initialized inner one. Each scope’s bindings are a **StrMap** (the same hash table described in §6.4), sized once per compilation from an exact upper bound — a single $O(n)$ AST walk counting every `NODE_VAR_DECL` before codegen starts — rather than a linked list scanned per lookup, for the same reason the optimizer passes need hashing rather than linear scanning: a wide, unnested program with many top-level declarations referenced later in original declaration order is precisely the access pattern that makes a linear scope scan degrade toward $O(v)$ per lookup (§6.6, Experiment B, exercises exactly this).

6.4 TAC-Level Optimization

`optimizeTAC()` runs four passes to a fixed point (capped at 8 iterations, since folding can expose new simplification opportunities and vice versa):

- **Constant folding** — a `BINOP/UNOP` whose operand(s) are literal integers is evaluated at compile time and rewritten to a plain `ASSIGN` of the computed literal (division by a literal zero is left unfolded, since that is a runtime error condition, not a compile-time constant).
- **Algebraic simplification** — identity/absorbing-element patterns (`x+0`, `0+x`, `x-0`, `x*1`, `1*x`, `x*0`, `0*x`, `x/1`) collapse a `BINOP` into a plain copy.
- **Local copy propagation** — a hash-table-backed alias map records, for each simple copy `Y = X`, both the source name `X` and a *version stamp*: the sequence number at which `X` was most recently (re)defined at the time of the copy. Every name gets its own monotonically-increasing sequence number, bumped on every `ASSIGN/BINOP/UNOP` result. A recorded copy is trusted — and chased through, transitively — only while the source’s *current* sequence number still matches the stamp taken at copy time; if `X` has been redefined since, the copy is treated as stale automatically, without any explicit scan. This makes invalidation an $O(1)$ side effect of bumping a definition’s sequence number, rather than a scan over every recorded alias. The alias map (but not the sequence-number history) is cleared at every label, since a label is a control-flow merge point where the straight-line assumption no longer holds.
- **Redundant-temporary / dead-code elimination** — since every generated temporary is consumed exactly once, a single linear scan collects every name still referenced anywhere in the surviving code into a hash set; any instruction defining a temporary that is not in that set is marked dead and dropped from further output. User-named variables are always kept (they are not subject to this pass), since a full whole-program liveness analysis across arbitrary control flow was judged out of scope for this project; this is a deliberate, documented simplifying assumption.

6.5 Complexity Analysis: Cost of AST Traversal and Code Emission

TAC generation walks the (already-built) AST exactly once, emitting a constant number of instructions per node (a `BINOP` emits one instruction; an `if/while` emits a small constant number of label/jump instructions in addition to its children’s code). For n AST nodes this is $O(n)$ time and produces $O(n)$ TAC instructions.

Each of the four optimization passes is a single linear scan over the instruction list. Achieving $O(m)$ per pass for m instructions specifically requires $O(1)$ -average lookup for both the copy-propagation alias map and the dead-code-elimination “used names” set — a flat array scanned linearly per lookup would instead give $O(m)$ *per lookup*, and $O(m^2)$ for the whole pass, since a pass performs up

to $O(m)$ lookups. Both maps are therefore backed by a chained hash table (djb2 hashing) rather than a flat array, giving $O(1)$ average-case lookup and insertion and restoring the intended $O(m)$ per-pass cost. With the fixed-point loop bounded by a constant (8) iteration cap, optimization is $O(m)$ overall (constant-factor, not growing with input) — confirmed empirically in §6.6, not just asserted.

6.6 Empirical Validation

Two properties are only as trustworthy as the evidence behind them: that the optimizer is genuinely $O(m)$ (§6.5), and that the §5.5 depth guard rejects oversized input instead of letting it run unboundedly or crash. Both are measured here against the actual `minicompiler` binary built from this repository, using two generator scripts and a harness committed at `testcases/bench/` (`gen_chain.py`, `gen_wide.py`, `bench.sh`) so every number below can be reproduced by running `testcases/bench/bench.sh` after `make` — not just read and trusted.

Experiment A — depth-guard behavior (`gen_chain.py`). Emits $x = 1+1+\dots+1$ with N terms: one deeply *nested* expression, exercising §5.5’s `MAX_EXPR_DEPTH = 3000` guard directly.

N (terms)	Result	Time
500 – 3,000	compiles	0.007 s – 0.021 s
3,001	rejected — expression nested too deeply (limit 3000 levels), exit 1	0.005 s
4,000 – 200,000	rejected, same diagnostic, exit 1	0.005 s – 0.067 s

The accept/reject boundary is exactly where §5.5 says it should be — nothing above 3,000 levels compiles, and every rejection is fast and clean rather than a hang or a crash, all the way out to 200,000 terms.

Experiment B — optimizer scaling (`gen_wide.py`). Emits N separate, shallow statements (`int x1; x1 = x0 + 1;`, repeated) instead of one nested expression, so each statement’s own expression depth stays $O(1)$ regardless of N — this isolates the optimizer’s per-instruction cost from the depth guard, letting m (TAC instruction count) grow into the hundreds of thousands without tripping §5.5.

N (statements)	Current (hash-table)	Legacy flat-array reference
2,000	0.029 s	1.643 s
4,000	0.057 s	5.968 s
8,000	0.118 s	28.729 s
16,000	0.247 s	did not finish in 60 s
32,000	0.523 s	—
50,000	0.808 s	—
100,000	1.650 s	—
200,000	3.350 s	—

The current column scales linearly across two orders of magnitude (100x the input gives ~115x

the time, not $\sim 10,000\times$). The legacy column is not a hypothetical: it is a real, reproducible measurement against `testcases/bench/legacy_optimizer_reference.c`, a preserved snapshot of `copyPropagatePass/deadCodeEliminatePass` as they existed before the §6.4/§6.5 hash-table rewrite (swap it into `codegen.c` in place of the current versions, per the instructions in that file, to reproduce this column directly). It shows textbook quadratic behavior — each doubling of N roughly quadruples the time ($1.6\text{ s} \rightarrow 6.0\text{ s} \rightarrow 28.7\text{ s}$) — and did not finish 16,000 statements within a 60-second timeout, where the current implementation takes 0.25 s. This is not a difference in constant factors; it is the difference between an algorithm that scales and one that does not, and it is exactly the class of bug this section exists to catch empirically rather than assume away theoretically.

The same investigation additionally found, and fixed, two related instances of the identical mistake — a component using a data structure whose cost degrades with scope size instead of staying $O(1)$ — that Experiment B alone would not have isolated from each other: `codegen.c`’s own scope-name resolver (§6.3) was a linked list scanned per lookup, and `symbol_table.c`’s per-scope hash table (§5.1) had a bucket count fixed at creation instead of growing with its entry count. Both are now backed by properly-sized or growable hash tables for the same reason as the optimizer passes, and both are exercised (not just theoretically implicated) by Experiment B, since it declares N distinct top-level variables and later references them in original declaration order — precisely the access pattern that would expose either bug.

Grammar note. Reaching $N = 200,000$ *statements* at all in Experiment B required one further fix, independent of TAC-optimizer cost: `parser.y`’s `stmt_list` rule was originally right-recursive, which forces Bison to hold every statement on its own parser stack until the whole file is parsed (§3.3) — a plain sequence of a few thousand statements failed to parse (`memory exhausted`) purely from grammar structure, before optimization was ever reached. It is now left-recursive with an $O(n)$ list-reversal pass, which is why Experiment B’s parser stack usage stays $O(1)$ per statement regardless of N .

The full test suite (10 valid programs exercising every language construct, 10 programs each triggering one required error class), plus both benchmark families above, were additionally run under AddressSanitizer and UndefinedBehaviorSanitizer (`gcc -fsanitize=address,undefined`) with zero reported memory-safety or undefined-behavior errors.

7. Overall Compiler Complexity

7.1 Combined Time Complexity of All Phases

Phase	Time Complexity	Dominant Cost
Lexical analysis	$O(n)$	n = source characters
Syntax analysis	$O(n)$	n = tokens (LALR table-driven)
AST construction	$O(n)$	one node per reduced construct
Semantic analysis	$O(n \cdot d)$	d = max scope-nesting depth (effectively $O(n)$)
TAC generation	$O(n)$	one AST walk

Phase	Time Complexity	Dominant Cost
TAC optimization	$O(m)$	$m = \text{TAC instructions}$, $O(n)$; hash-table-backed passes, constant-bounded iteration
Target code generation	$O(m)$	one TAC walk

Since every phase is linear (or effectively linear, treating scope-nesting depth as a small constant bounded by realistic program structure) in the size of its input, and each phase’s input size is itself $O(n)$ in the original source length, the compiler as a whole runs in **$O(n)$** time for a source program of n characters. §6.6 confirms this empirically, not just theoretically, for the phase (TAC optimization) where the implementation originally diverged from the claim.

7.2 Space Usage of AST, Symbol Table, and TAC

- **AST:** $O(n)$ nodes, each of fixed size — $O(n)$ total.
- **Symbol table:** $O(v)$ live symbols at any point in the walk, where v is the number of variables declared in currently-open scopes; peak total across the whole compilation is $O(V)$ where V is the total number of declarations in the program.
- **TAC:** $O(n)$ instructions (bounded by the AST size that generated them), stored in a dynamic array that doubles on growth, giving $O(1)$ amortized append and $O(n)$ total space. The optimizer’s hash tables (§6.4) are sized proportionally to the instruction count per pass invocation and are discarded at the end of each pass, adding $O(m)$ space, not $O(m)$ *per iteration accumulated*.

7.3 Discussion: Does the Compiler Operate in Linear Time?

Yes, for realistic MiniLang programs, and this is no longer solely a theoretical claim — §6.6 measured it directly across two orders of magnitude of input size and confirmed linear scaling, and confirmed that the case that previously broke this guarantee (the optimizer’s original flat-array implementation) has been fixed and re-measured. Every phase is either strictly $O(n)$ (lexing, parsing, AST construction, TAC generation, target codegen) or $O(n \cdot d)$ with d bounded by scope-nesting depth (semantic analysis) — and d grows with how deeply the programmer chooses to nest `{ }` blocks, not with program length, so it is effectively a small constant in practice. The TAC-optimization fixed-point loop is capped at a constant number of iterations rather than being allowed to scale with program size, which keeps it from becoming a hidden superlinear cost. Bounded recursion depth (§5.5) additionally guarantees that inputs which *would* exceed a safe C-stack depth are rejected in $O(1)$ time with a clean diagnostic rather than being allowed to run unboundedly or crash. There is no phase in this implementation that performs quadratic or worse work as a function of source size, so the compiler is linear end-to-end.

8. Sample Compilation Results

8.1 Example MiniLang Program (`testcases/valid1_features.ml`)

```
// Exercises every supported MiniLang construct.
int a;
int b;
```

```

int sum;
bool flag;

a = 5;
b = 10;
sum = a + b * 2;
flag = sum > 20;

if (flag) {
    print(sum);
} else {
    print(0);
}

int i;
i = 0;
while (i < 3) {
    print(i);
    i = i + 1;
}

```

8.2 Generated Three-Address Code (unoptimized)

```

a = 5
b = 10
$t1 = b * 2
$t2 = a + $t1
sum = $t2
$t3 = sum > 20
flag = $t3
ifz flag goto $L1
print sum
goto $L2
$L1:
print 0
$L2:
i = 0
$L3:
$t4 = i < 3
ifz $t4 goto $L4
print i
$t5 = i + 1
i = $t5
goto $L3
$L4:

```

8.3 Optimized Three-Address Code (written to `output.tac`)

```
a = 5
b = 10
sum = 25
flag = 1
ifz 1 goto $L1
print 25
goto $L2
$L1:
print 0
$L2:
i = 0
$L3:
$t4 = i < 3
ifz $t4 goto $L4
print i
$t5 = i + 1
i = $t5
goto $L3
$L4:
```

Constant folding collapses `b * 2` (10×2) and `a + $t1` ($5 + 20$) directly to `sum = 25`, and `sum > 20` folds to `flag = 1`; the now-constant branch condition and the loop body (whose trip count depends on a runtime variable, `i`) are left as genuine control flow, as required. This output was re-verified byte-for-byte against the current implementation after the §6.4–6.6 optimizer rewrite, confirming the rewrite changed the optimizer’s internal data structures without changing its observable behavior on any existing test case.

8.4 Generated Target (Stack-Machine) Code (written to `output.asm`)

```
LOAD 5      STORE a      LOAD 10     STORE b
LOAD 25     STORE sum    LOAD 1      STORE flag
LOAD 1      JZ $L1       LOAD 25     PRINT
JMP $L2     $L1: LOAD 0 PRINT
$L2: LOAD 0 STORE i      $L3:
LOAD i      LOAD 3       LT          STORE $t4
LOAD $t4    JZ $L4       LOAD i      PRINT
LOAD i      LOAD 1       ADD         STORE $t5
LOAD $t5    STORE i      JMP $L3     $L4:
```

(Reflowed onto four columns here to fit the page; the actual `output.asm` has one instruction per line.) Every TAC instruction maps mechanically to LOAD/STORE around one operation, e.g. `t = a op b` $\rightarrow$ `LOAD a; LOAD b; OP; STORE t` — which sidesteps register allocation entirely by treating every temporary and variable as an addressable storage slot.

8.5 Verified Error Handling

Each error class was exercised against the actual built `minicompiler` and produced the expected diagnostic with a correct exit code of 1. Testcase names below are relative to `testcases/` and,

except the last row, end in .ml:

Testcase	Type	Sample Output
err_lexical	Lexical	Lexical Error at line 2: unexpected character '@'
err_lexical_terminated_comment	Lexical	Lexical Error at line 4: terminated comment
err_syntax	Syntax	Syntax Error at line 2: syntax error, unexpected IDENT, expecting ASSIGN
err_semantic_undeclared	Semantic	Semantic Error at line 2: undeclared variable 'b'
err_semantic_redeclare	Semantic	Semantic Error at line 2: variable 'a' is already declared in this scope
err_semantic_typemismatch	Semantic	Semantic Error at line 3: type mismatch: cannot assign bool value to int variable 'a'
err_semantic_condition	Semantic	Semantic Error at line 6: invalid conditional expression: expected bool, got int
err_semantic_binop_typemismatch	Semantic	Semantic Error at line 6: type mismatch: operator '+' requires int operands, got bool and bool
err_semantic_eq_typemismatch	Semantic	Semantic Error at line 5: type mismatch: operands of '==' must have the same type, got int and bool
err_semantic_unary_typemismatch	Semantic	Semantic Error at line 4: type mismatch: unary '-' requires an int operand, got bool
(synthetic stress input)	Robustness (§5.5)	Semantic Error at line 2: expression nested too deeply (limit 3000 levels)

8.6 Verified Scope Shadowing (testcases/valid2_scoping.ml)

This test specifically exercises the codegen-level scope handling described in §6.3:

```
int x; x = 1;
{ int x; x = 99; print(x); }
print(x);
```

TAC	Optimized
x = 1	x = 1
x\$2 = 99	x\$2 = 99
print x\$2	print 99
print x	print 1

The inner block's `x` is correctly disambiguated to `x$2`, so the two `print` statements resolve to their

own distinct variables and correctly print 99 then 1, matching MiniLang's block-scoping semantics.